

```
Data_Analysis_and_Regression (/github/eesh400/Data_Analysis_and_Regression/tree/main)
/
Data_Analysis_and_Regression.ipynb (/github/eesh400/Data_Analysis_and_Regression/tree/main/Data_Analysis_and_Regression.ipynb)
```

```
In [1]: #-----Exploring different types of data(N,O,I,R)-----
import pandas as pd
df=pd.DataFrame({'City': ['Delhi', 'Mumbai', 'Kolkata', 'Delhi'],           #Nominal
                 'Satisfaction':[1,2,3,5],      #Ordinal
                 'Temperature': [22.5, 27.4, 28, 14],    #Interval
                 'Income': [20000, 45000, 68900, 56000]})   #Ratio

print(df.dtypes)          #If one or more entries are missing, they are shown as NaN
print(df.describe(include='all'))      #Missing values as NaN, use df['Satisfaction'].dropna() to
```

```
City          object
Satisfaction   int64
Temperature    float64
Income         int64
dtype: object
      City  Satisfaction  Temperature        Income
count      4        4.000000     4.000000    4.000000
unique     3            NaN            NaN            NaN
top       Delhi          NaN          NaN          NaN
freq       2            NaN          NaN          NaN
mean      NaN        2.750000    22.975000  47475.000000
std       NaN        1.707825    6.470639  20758.191154
min       NaN        1.000000    14.000000  20000.000000
25%      NaN        1.750000    20.375000  38750.000000
50%      NaN        2.500000    24.950000  50500.000000
75%      NaN        3.500000    27.550000  59225.000000
max       NaN        5.000000    28.000000  68900.000000
```

```
In [15... # Handling Missing and dirty data
import numpy as np
import pandas as pd
df= pd.DataFrame({'age':[22, 25, np.nan, 30, 28, 34],
                  'salary':[30000, 50000, 45000, np.nan, 60000, 52000]})

# 1. Check missing data
print("Missing Data:\n", df.isnull().sum())

# 2. Imputation
df['age']= df['age'].fillna(df['age'].mean())      #mean imputation
df['salary']= df['salary'].fillna(df['salary'].median())    #median imputation
print("After imputation: \n", df)

# 3. Outlier Detection
q1=df['salary'].quantile(0.25)
q3=df['salary'].quantile(0.75)
iqr= q3-q1

lower=q1 - 1.5*iqr
upper=q3 + 1.5*iqr

outliers=df[(df['salary']< lower)
             |(df['salary']>upper)]
print("Outliers are:\n", outliers)

# 3. Deletion of missing data
df_drop_salary= df.dropna(subset=['salary']) #Rows with missing values in specific columns deleted
print("After deletion:\n",df_drop_salary)      #not deleted since values already assigned
#pandas can't delete single cell, removes rows or columns
# use df_drop=df.dropna() print("After deletion:", df_drop)
```

Missing Data:

|   | age  | salary  |
|---|------|---------|
| 0 | 22.0 | 30000.0 |
| 1 | 25.0 | 50000.0 |
| 2 | 27.8 | 45000.0 |
| 3 | 30.0 | 50000.0 |
| 4 | 28.0 | 60000.0 |
| 5 | 34.0 | 52000.0 |

After imputation:

|   | age  | salary  |
|---|------|---------|
| 0 | 22.0 | 30000.0 |
| 1 | 25.0 | 50000.0 |
| 2 | 27.8 | 45000.0 |
| 3 | 30.0 | 50000.0 |
| 4 | 28.0 | 60000.0 |
| 5 | 34.0 | 52000.0 |

Outliers are:

|   | age  | salary  |
|---|------|---------|
| 0 | 22.0 | 30000.0 |
| 4 | 28.0 | 60000.0 |

After deletion:

|   | age  | salary  |
|---|------|---------|
| 0 | 22.0 | 30000.0 |
| 1 | 25.0 | 50000.0 |
| 2 | 27.8 | 45000.0 |
| 3 | 30.0 | 50000.0 |
| 4 | 28.0 | 60000.0 |
| 5 | 34.0 | 52000.0 |

```
In [11... # Min-Max Normalization(0 to 1)

from sklearn.preprocessing import MinMaxScaler
scaler= MinMaxScaler()
df[['age_norm', 'salary_norm']] = scaler.fit_transform(df[['age', 'salary']])
print("Min-Max normalized:\n", df)
```

Min-Max normalized:

|   | age  | salary  | age_norm | salary_norm |
|---|------|---------|----------|-------------|
| 0 | 22.0 | 30000.0 | 0.000000 | 0.000000    |
| 1 | 25.0 | 50000.0 | 0.250000 | 0.666667    |
| 2 | 27.8 | 45000.0 | 0.483333 | 0.500000    |
| 3 | 30.0 | 50000.0 | 0.666667 | 0.666667    |
| 4 | 28.0 | 60000.0 | 0.500000 | 1.000000    |
| 5 | 34.0 | 52000.0 | 1.000000 | 0.733333    |

```
In [13...]: # Standardization (Mean=0, Std.=1)
from sklearn.preprocessing import StandardScaler
std= StandardScaler()
df[['age_std', 'salary_std']] = std.fit_transform(df[['age', 'salary']])

print("Standardized:\n", df)
```

Standardized:

|   | age  | salary  | age_norm | salary_norm | age_std   | salary_std |
|---|------|---------|----------|-------------|-----------|------------|
| 0 | 22.0 | 30000.0 | 0.000000 | 0.000000    | -1.542786 | -1.951918  |
| 1 | 25.0 | 50000.0 | 0.250000 | 0.666667    | -0.744793 | 0.237149   |
| 2 | 27.8 | 45000.0 | 0.483333 | 0.500000    | 0.000000  | -0.310118  |
| 3 | 30.0 | 50000.0 | 0.666667 | 0.666667    | 0.585195  | 0.237149   |
| 4 | 28.0 | 60000.0 | 0.500000 | 1.000000    | 0.053200  | 1.331682   |
| 5 | 34.0 | 52000.0 | 1.000000 | 0.733333    | 1.649185  | 0.456056   |

```
In [4]: import pandas as pd
import numpy as np
df1= pd.DataFrame({ 'Circlename':["Telangana Circle", "Telangana Circle", "Telangana Circle", "Telangana Circle"],
                    'Regionname':["Hyderabad", "Hyderabad", "Hyderabad", "Hyderabad"],
                    'Officename':["Kothimir B.O", "Papanpet B.O", "Kukuda B.O", "Bareguda B.O"],
                    'Pincode': [504273, 504299, 504299, 504296],
                    'Latitude':[19.3638689, 19.4764899, np.nan, 19.3285752],
                    'Population': [15000, 6708, 32000, 12574]})

#1. Basic data analysis
print("A. View data:\n", df1.head())      #view data
print("B. Summary:\n", df1.describe())    # summary stats
print("C. Correlation:\n", df1.corr(numeric_only=True))    # correlation
```

A. View data:

|   | Circlename       | Regionname | Officename   | Pincode | Latitude  | Population |
|---|------------------|------------|--------------|---------|-----------|------------|
| 0 | Telangana Circle | Hyderabad  | Kothimir B.O | 504273  | 19.363869 | 15000      |
| 1 | Telangana Circle | Hyderabad  | Papanpet B.O | 504299  | 19.476490 | 6708       |
| 2 | Telangana Circle | Hyderabad  | Kukuda B.O   | 504299  | NaN       | 32000      |
| 3 | Telangana Circle | Hyderabad  | Bareguda B.O | 504296  | 19.328575 | 12574      |

B. Summary:

|       | Pincode       | Latitude  | Population   |
|-------|---------------|-----------|--------------|
| count | 4.000000      | 3.000000  | 4.000000     |
| mean  | 504291.750000 | 19.389645 | 16570.500000 |
| std   | 12.579746     | 0.077253  | 10859.356319 |
| min   | 504273.000000 | 19.328575 | 6708.000000  |
| 25%   | 504290.250000 | 19.346222 | 11107.500000 |
| 50%   | 504297.500000 | 19.363869 | 13787.000000 |
| 75%   | 504299.000000 | 19.420179 | 19250.000000 |
| max   | 504299.000000 | 19.476490 | 32000.000000 |

C. Correlation:

|            | Pincode  | Latitude  | Population |
|------------|----------|-----------|------------|
| Pincode    | 1.000000 | 0.388297  | 0.128891   |
| Latitude   | 0.388297 | 1.000000  | -0.868328  |
| Population | 0.128891 | -0.868328 | 1.000000   |

In [6]: #Simple Linear Regression

```
df1['Latitude']= pd. to_numeric(df1['Latitude'], errors='coerce')
df1['Latitude']= df1['Latitude'].fillna(df1['Latitude'].mean())

X = df1[['Pincode']]
y =df1['Latitude']

from sklearn.linear_model import LinearRegression
import numpy as np

model= LinearRegression()
model.fit(X, y)

print("Coefficient:", model.coef_)
print("Intercept:", model.intercept_)
print("Predictions:", model.predict(X))

Coefficient: [0.00179753]
Intercept: -887.0909306440465
Predictions: [19.35594094 19.40267677 19.40267677 19.39728418]
```

In [7]: # Multiple LR

```
#Predict Latitude from pincode and population
X_multi = df1[['Pincode', 'Population']]
y1= df1['Latitude']

multi= LinearRegression()
multi.fit(X_multi,y1)

print("Coefficients:", multi.coef_)
print("Intercept:", multi.intercept_)

Coefficients: [ 2.01082245e-03 -1.91697902e-06]
Intercept: -994.6197618924509
```

In [8]:

```
import numpy as np
result=[5,4,3,4,5,3,5,4,2,3,4,2,3,1,5,2,1,4,1,2,4]
temp=[31,32,45,12,23,45]
sorted_result=sorted(result)
print(sorted_result)
print(np.mean(result))
print(np.median(result))
print(np.mean(temp))
```

```
[1, 1, 1, 2, 2, 2, 3, 3, 3, 4, 4, 4, 4, 4, 4, 5, 5, 5, 5]
3.1904761904761907
3.0
31.33333333333332
```

In [9]:

```
squared_differences=[]
for i in temp:
    diff= i-np.mean(temp)
    squared_difference=diff**2
    squared_differences.append(squared_difference)
avg_sqd=np.mean(squared_differences)
standard_deviation=np.sqrt(avg_sqd)
print(standard_deviation)
```

```
13.666666666666668
```

```
In [10... #geometric mean
num_items=len(temp)
product=1

for i in temp:
    product*=i
geo_mean=product**(1/num_items)
print(geo_mean)

28.662075145783213
```

```
In [11... import pandas as pd

list=[1,2,3,4,5,6,7,8]
series=pd.Series(list)
print(series)
```

```
0    1
1    2
2    3
3    4
4    5
5    6
6    7
7    8
dtype: int64
```

```
In [14... dict1={'state':['Assam', 'Delhi', 'Kerala'], 'GArea':[76589, 56794, 75436], 'VADF':[3657, 5735, 6863]
print(dict1)
dframe=pd.DataFrame(dict1)
print(dframe)
```

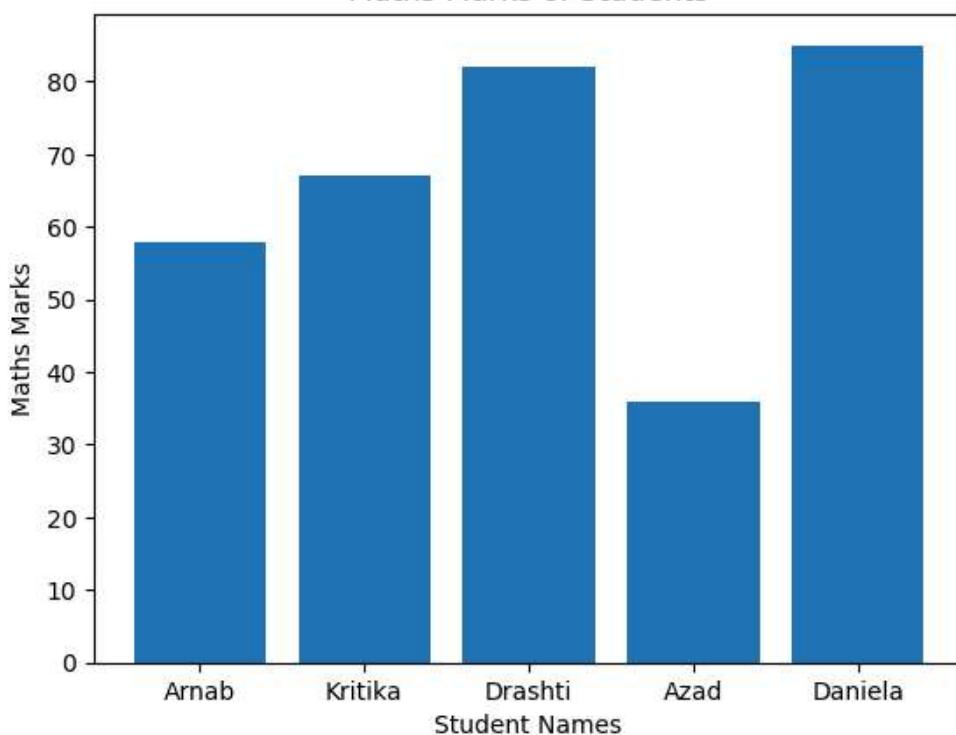
```
{'state': ['Assam', 'Delhi', 'Kerala'], 'GArea': [76589, 56794, 75436], 'VADF': [3657, 5735, 6863]}
   state   GArea   VADF
0  Assam    76589   3657
1   Delhi    56794   5735
2  Kerala    75436   6863
```

```
In [21... #2nd table
marks=pd.DataFrame([[1, 'Arnab', 16, 58], [2, 'Kritika', 32, 67], [3, 'Drashti', 24, 82],[4, 'Azad',
marks['Pass']=(marks['Maths']>45)
marks
```

|          | Roll No. | Name    | Eco | Maths | Pass  |
|----------|----------|---------|-----|-------|-------|
| <b>0</b> | 1        | Arnab   | 16  | 58    | True  |
| <b>1</b> | 2        | Kritika | 32  | 67    | True  |
| <b>2</b> | 3        | Drashti | 24  | 82    | True  |
| <b>3</b> | 4        | Azad    | 67  | 36    | False |
| <b>4</b> | 5        | Daniela | 57  | 85    | True  |

```
In [22... import matplotlib.pyplot as plt
plt.bar(marks['Name'], marks['Maths'])
plt.xlabel("Student Names")
plt.ylabel("Maths Marks")
plt.title("Maths Marks of Students")
plt.show()
```

### Maths Marks of Students



```
In [28...]: import pandas as pd
data={0:[1, 2, 3, 4, 5], 'TV':[230.1, 44.5, 17.2, 151.5, 180.8], 'Radio':[37.8, 39.3, 45.9, 41.3, 10.8]
df2=pd.DataFrame(data)
df2.to_csv("Advertising.csv", index=False)
df2=pd.read_csv("Advertising.csv")
print(df2.head())
#del column
df2.drop(df2.columns[0], axis=1, inplace=True)
df2.corr()
```

|   | 0 | TV    | Radio | News | Sales |
|---|---|-------|-------|------|-------|
| 0 | 1 | 230.1 | 37.8  | 69.2 | 22.2  |
| 1 | 2 | 44.5  | 39.3  | 45.1 | 10.4  |
| 2 | 3 | 17.2  | 45.9  | 69.3 | 9.3   |
| 3 | 4 | 151.5 | 41.3  | 58.5 | 18.5  |
| 4 | 5 | 180.8 | 10.8  | 58.4 | 12.9  |

```
Out[28...]:
```

|              | TV        | Radio     | News     | Sales    |
|--------------|-----------|-----------|----------|----------|
| <b>TV</b>    | 1.000000  | -0.478276 | 0.285756 | 0.857584 |
| <b>Radio</b> | -0.478276 | 1.000000  | 0.166737 | 0.036445 |
| <b>News</b>  | 0.285756  | 0.166737  | 1.000000 | 0.363963 |
| <b>Sales</b> | 0.857584  | 0.036445  | 0.363963 | 1.000000 |

```
In [29...]: #LR
from sklearn.linear_model import LinearRegression
x1=df2[["TV"]]
y1=df2[["Sales"]]
print(x1.head())
model=LinearRegression()
model.fit(x1,y1)
```

```
TV
0 230.1
1 44.5
2 17.2
3 151.5
4 180.8
```

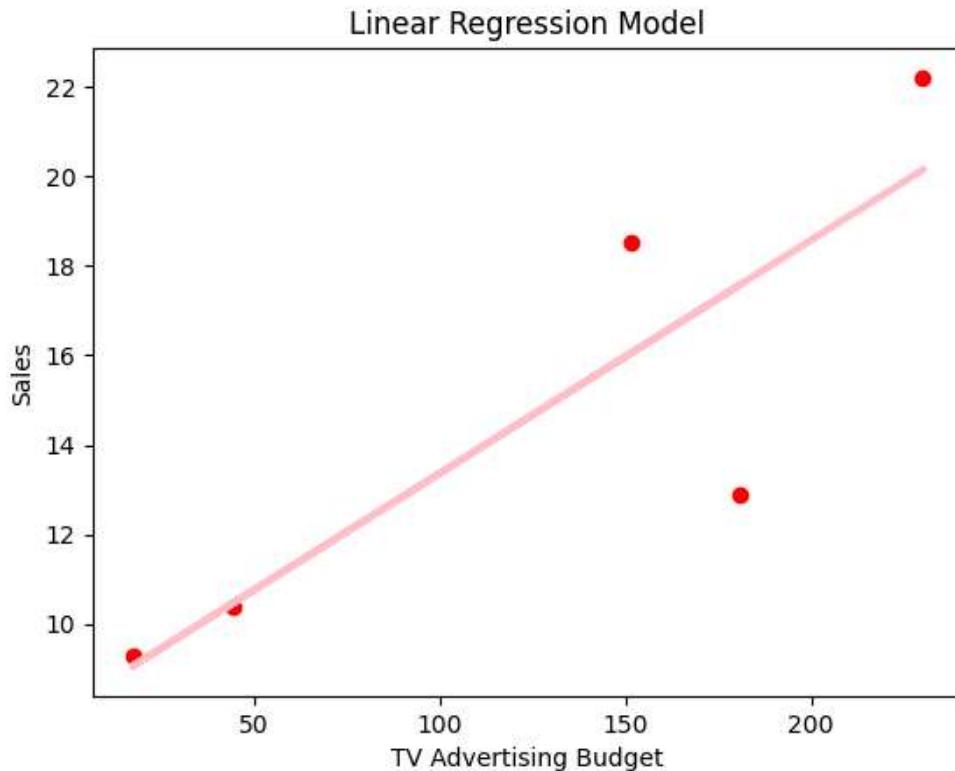
Out[29...]: LinearRegression

i ? ([https://scikit-learn.org/1.6/modules/generated/sklearn.linear\\_model.LinearRegression.html](https://scikit-learn.org/1.6/modules/generated/sklearn.linear_model.LinearRegression.html))

```
LinearRegression()
```

In [31...]:

```
a_1=model.coef_
a_0=model.intercept_
plt.scatter(x1,y1,color='red')
plt.plot(x1,a_0+a_1*x1, color='pink', linewidth=3)
plt.title('Linear Regression Model')
plt.xlabel('TV Advertising Budget')
plt.ylabel('Sales')
plt.show()
```



In [33...]:

```
print(model.coef_)
print(a_0)
```

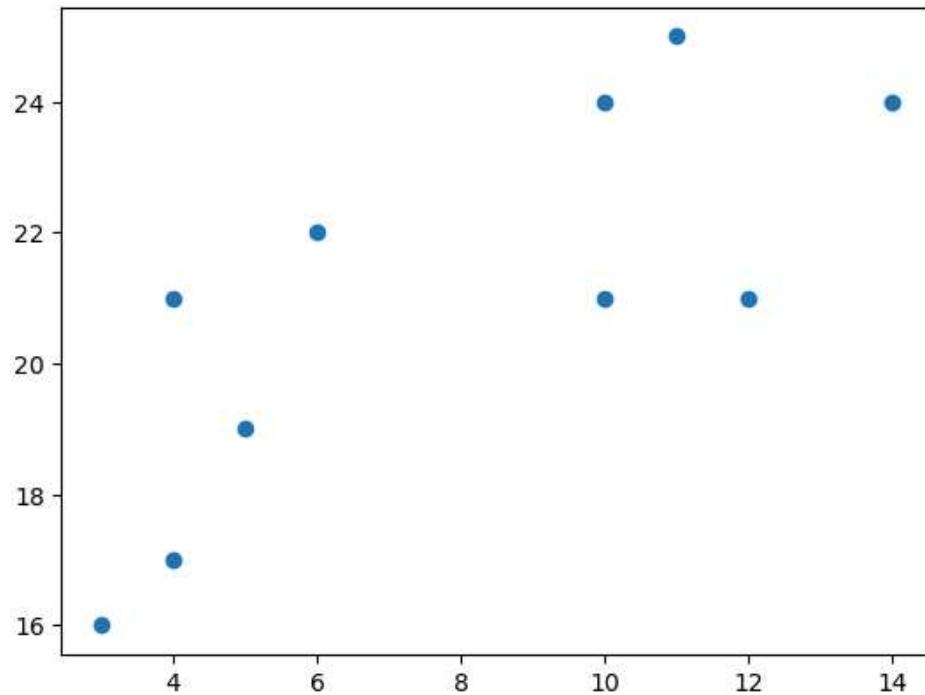
```
[[0.05208121]]
[8.15922324]
```

In [2]:

```
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans

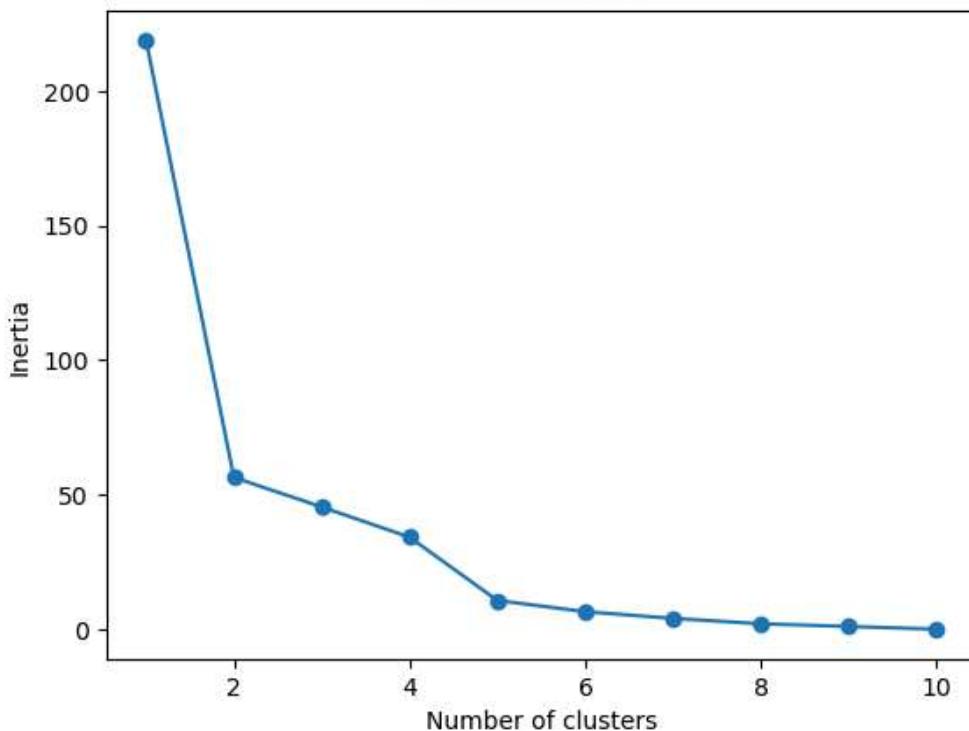
x3 = [4, 5, 10, 4, 3, 11, 14, 6, 10, 12]
y3 = [21, 19, 24, 17, 16, 25, 24, 22, 21, 21]

plt.scatter(x3, y3) #Scatter plot visualizes the data distribution prior to applying K-means
plt.show()
```



```
In [4]: from sklearn.cluster import KMeans  
  
data3 = list(zip(x3, y3)) # Using the elbow method to determine the optimal number of clusters  
inertias = []  
  
for i in range(1,11):  
    kmeans = KMeans(n_clusters=i)  
    kmeans.fit(data3)  
    inertias.append(kmeans.inertia_)  
  
plt.plot(range(1,11), inertias, marker='o')  
plt.title('Elbow method')  
plt.xlabel('Number of clusters')  
plt.ylabel('Inertia')  
plt.show()
```

## Elbow method



```
In [40... del list
```

```
In [44... from sklearn.linear_model import LinearRegression  
  
x4= df2[['TV', 'Radio']]  
y4= df2['Sales']  
multi4= LinearRegression()  
multi4.fit(x4,y4)  
  
print("Coefficients:", multi.coef_)  
print("Intercept:", multi.intercept_)
```

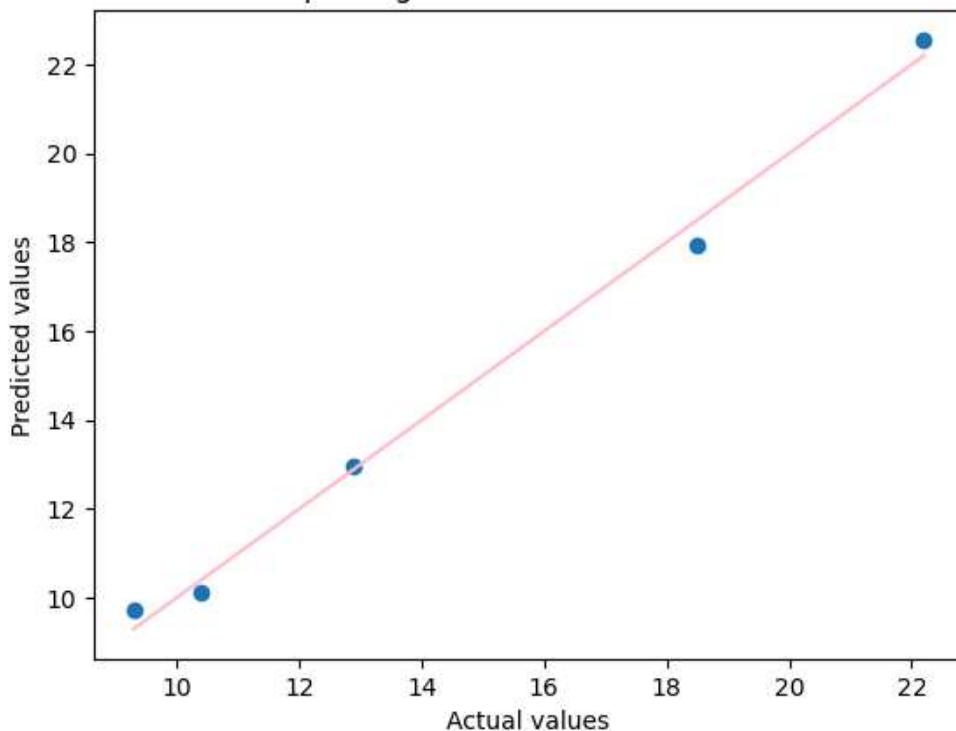
```
Coefficients: [ 2.01082245e-03 -1.91697902e-06]  
Intercept: -994.6197618924509
```

```
In [49... predicted_Sales=multi4.predict(pd.DataFrame([[100,150]],columns=['TV','Radio']))  
print(predicted_Sales)
```

```
[39.39772356]
```

```
In [55... import matplotlib.pyplot as plt  
y_pred=multi4.predict(x4)  
plt.scatter(y4, y_pred)  
plt.plot([y4.min(), y4.max()],[y4.min(),y4.max()], color='pink')  
plt.xlabel("Actual values")  
plt.ylabel("Predicted values")  
plt.title("Multiple Regression: Actual vs. Predicted")  
plt.show()
```

### Multiple Regression: Actual vs. Predicted



```
In [5]: import matplotlib.pyplot as plt
from scipy import stats

x5 = [5,7,8,7,2,17,2,9,4,11,12,9,6]
y5 = [99,86,87,88,111,86,103,87,94,78,77,85,86]

slope, intercept, r, p, std_err = stats.linregress(x5, y5)

def myfunc(x):
    return slope * x + intercept

mymodel = list(map(myfunc, x5))                      #Estimating a Linear relationship between x5 and y5

plt.scatter(x5, y5)
plt.plot(x5, mymodel)
plt.show()
```

